

# Task

## Iris Flower Classification

Iris flower has three species; setosa, versicolor, and virginica, which differs according to their measurements. Now assume that you have the measurements of the iris flowers according to their species, and here your task is to train a machine learning model that can learn from the measurements of the iris species and classify them. Although the Scikit-learn library provides a dataset for iris flower classification.

## Import Librarys

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
```

## Load dataset

```
df = pd.read_csv("Iris.csv")
```

```
df.head()
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	
Species						
0	1	5.1	3.5	1.4	0.2	Iris-
						setosa
1	2	4.9	3.0	1.4	0.2	Iris-
						setosa
2	3	4.7	3.2	1.3	0.2	Iris-
						setosa
3	4	4.6	3.1	1.5	0.2	Iris-
						setosa
4	5	5.0	3.6	1.4	0.2	Iris-
						setosa

# Data Information

```
df.shape
```

```
(150, 5)
```

```
df.isnull().sum()
```

```
Id          0
SepalLengthCm  0
SepalWidthCm  0
PetalLengthCm  0
PetalWidthCm  0
Species      0
dtype: int64
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 150 entries, 0 to 149
```

```
Data columns (total 6 columns):
```

#	Column	Non-Null Count	Dtype
0	Id	150 non-null	int64
1	SepalLengthCm	150 non-null	float64
2	SepalWidthCm	150 non-null	float64
3	PetalLengthCm	150 non-null	float64
4	PetalWidthCm	150 non-null	float64
5	Species	150 non-null	object

```
dtypes: float64(4), int64(1), object(1)
```

```
memory usage: 7.2+ KB
```

```
df.columns
```

```
Index(['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm',  
      'PetalWidthCm',  
      'Species'],  
      dtype='object')
```

```
df.nunique()
```

```
Id          150
SepalLengthCm  35
SepalWidthCm  23
PetalLengthCm  43
PetalWidthCm  22
Species       3
dtype: int64
```

```
df['Species'].unique()
```

```
array(['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'],  
      dtype=object)
```

## Drop 'Id' column as it's not needed

```
df.drop(columns=['Id'], inplace=True)  
  
df.columns  
Index(['SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm',  
      'PetalWidthCm',  
      'Species'],  
      dtype='object')
```

## Encode species labels

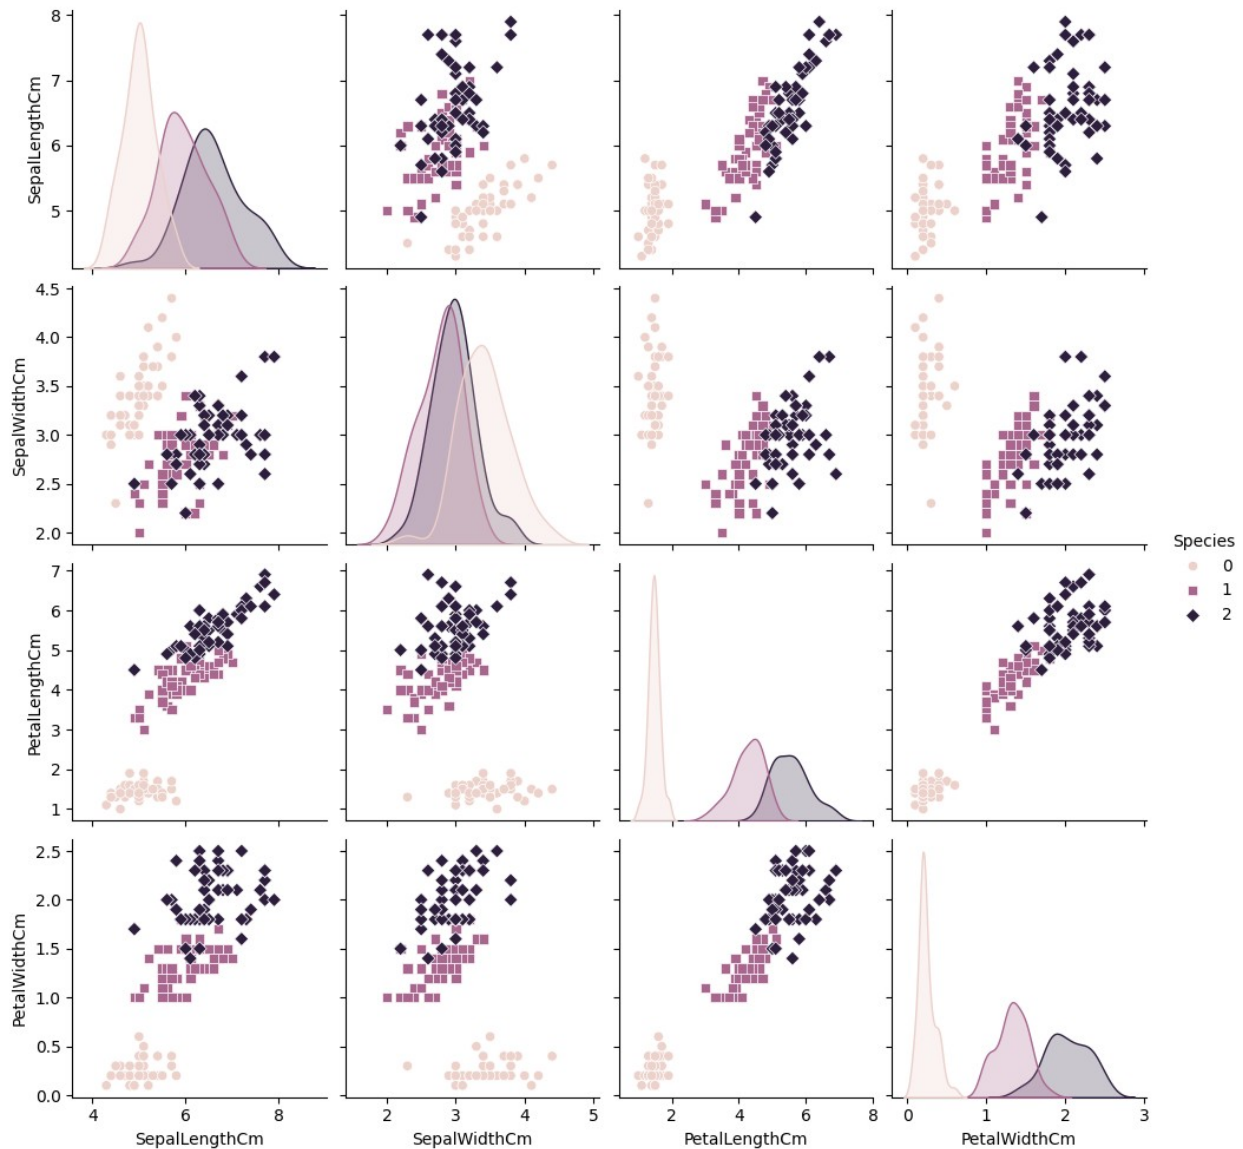
```
label_encoder = LabelEncoder()  
df['Species'] = label_encoder.fit_transform(df['Species'])  
  
df.head()  


|   | SepalLengthCm | SepalWidthCm | PetalLengthCm | PetalWidthCm | Species |
|---|---------------|--------------|---------------|--------------|---------|
| 0 | 5.1           | 3.5          | 1.4           | 0.2          | 0       |
| 1 | 4.9           | 3.0          | 1.4           | 0.2          | 0       |
| 2 | 4.7           | 3.2          | 1.3           | 0.2          | 0       |
| 3 | 4.6           | 3.1          | 1.5           | 0.2          | 0       |
| 4 | 5.0           | 3.6          | 1.4           | 0.2          | 0       |

  
df['Species'].unique()  
array([0, 1, 2])
```

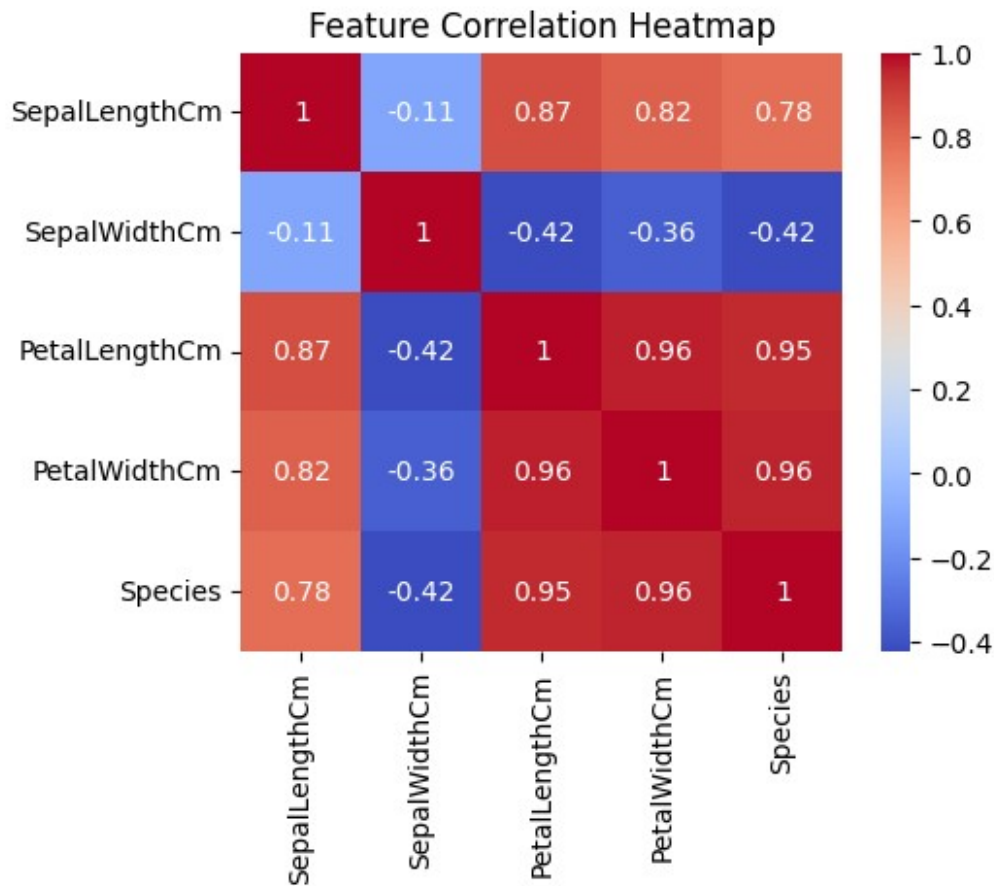
## --- Data Visualization ---

```
sns.pairplot(df, hue="Species", markers=["o", "s", "D"])  
plt.show()
```



*# Pairplot shows relationships between features for different species.*

```
plt.figure(figsize=(5, 4))
sns.heatmap(df.corr(), annot=True, cmap="coolwarm")
plt.title("Feature Correlation Heatmap")
plt.show()
```



```
# Heatmap shows correlations between sepal/petal measurements.
```

## Split features and target

```
X = df.drop(columns=['Species'])  
y = df['Species']
```

## Split data into training and test sets (80-20 split)

```
X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.2, random_state=42)
```

## Train the model (Random Forest)

```
model = RandomForestClassifier(n_estimators=100, random_state=42)
model.fit(X_train, y_train)
```

```
RandomForestClassifier(random_state=42)
```

## Make predictions

```
y_pred = model.predict(X_test)
```

```
y_pred
```

```
array([1, 0, 2, 1, 1, 0, 1, 2, 1, 1, 2, 0, 0, 0, 0, 1, 2, 1, 1, 2, 0,
       2,
       0, 2, 2, 2, 2, 2, 0, 0])
```

## Evaluate the model

```
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
print("Classification Report:\n", classification_report(y_test,
y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11

accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

Confusion Matrix:

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```

# Train Machine Learning Models (SVM, k-NN, Decision Tree)

```
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
```

## Train models

```
models = {
    "SVM": SVC(kernel='linear'),
    "k-NN": KNeighborsClassifier(n_neighbors=5),
    "Decision Tree": DecisionTreeClassifier(random_state=42)
}
```

## Evaluate models

```
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    print(f"\n=== {name} Model Performance ===")
    print("Accuracy:", accuracy_score(y_test, y_pred))
    print("Classification Report:\n", classification_report(y_test,
y_pred))
    print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

=== SVM Model Performance ===

Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

Confusion Matrix:

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```

=== k-NN Model Performance ===

Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

Confusion Matrix:

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```

=== Decision Tree Model Performance ===

Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

Confusion Matrix:

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```

## Select only two features for visualization

```
X = df[['PetalLengthCm', 'PetalWidthCm']] # Use only PetalLength and PetalWidth
```

```
y = df["Species"]
```

```
# Train-test split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

```
# Define models and train them on the selected features
```

```
models = {
```



```

    "SVM": SVC(kernel='linear'),
    "k-NN": KNeighborsClassifier(n_neighbors=5),
    "Decision Tree": DecisionTreeClassifier(random_state=42),
    "Random Forest": RandomForestClassifier(n_estimators=10,
random_state=42)
}

for name, model in models.items():
    model.fit(X_train, y_train) # Train using only the two selected
features

def plot_decision_boundary(model, X, y, title):
    x_min, x_max = X.iloc[:, 0].min() - 1, X.iloc[:, 0].max() + 1
    y_min, y_max = X.iloc[:, 1].min() - 1, X.iloc[:, 1].max() + 1
    xx, yy = np.meshgrid(np.linspace(x_min, x_max, 200),
np.linspace(y_min, y_max, 200))

    # Predict labels for each point in mesh grid
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)


    # Plot decision boundary
    plt.contourf(xx, yy, Z, alpha=0.3, cmap=plt.cm.coolwarm)
    sns.scatterplot(x=X.iloc[:, 0], y=X.iloc[:, 1], hue=y,
palette="Dark2", edgecolor="k")

    plt.xlabel(X.columns[0])
    plt.ylabel(X.columns[1])
    plt.title(title)
    plt.show()

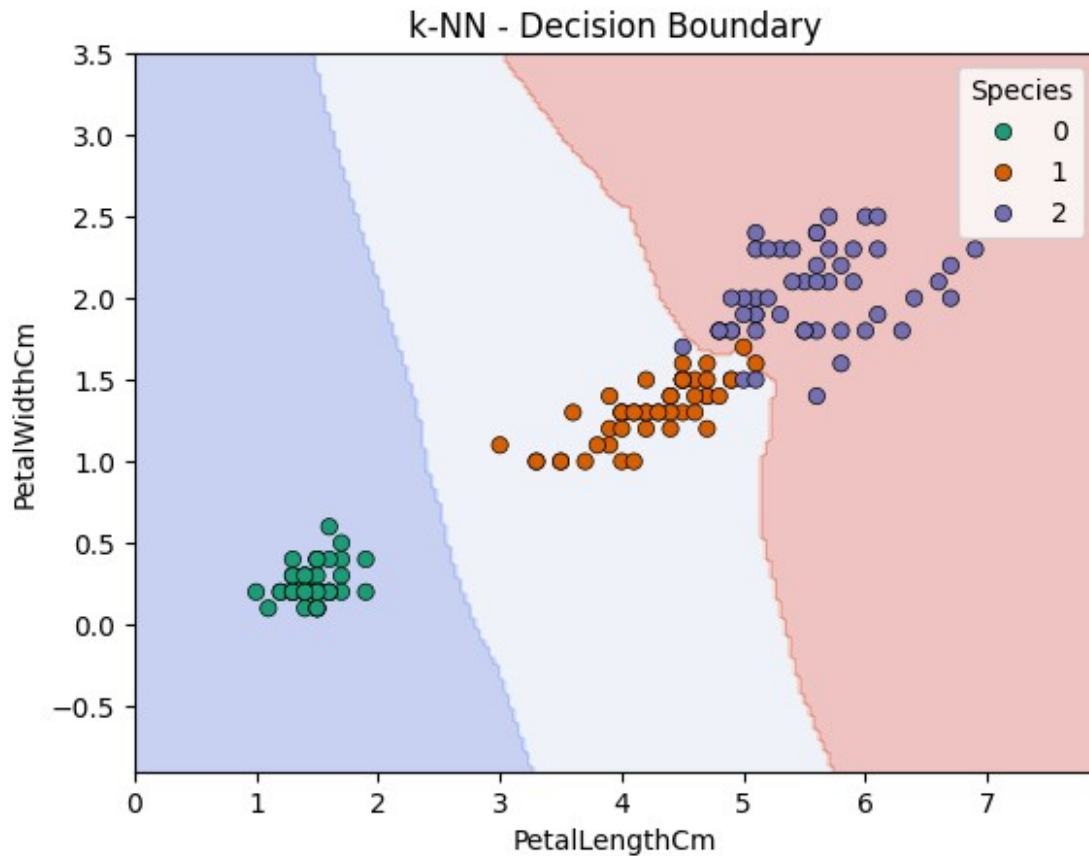
# Plot decision boundaries for each model
for name, model in models.items():
    plot_decision_boundary(model, X, y, f"{name} - Decision Boundary")

C:\Users\Admin\AppData\Local\Programs\Python\Python313\Lib\site-
packages\sklearn\utils\validation.py:2739: UserWarning: X does not
have valid feature names, but SVC was fitted with feature names
warnings.warn(


```




```
C:\Users\Admin\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\utils\validation.py:2739: UserWarning: X does not have valid feature names, but KNeighborsClassifier was fitted with feature names
  warnings.warn(
```



```
C:\Users\Admin\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\utils\validation.py:2739: UserWarning: X does not have valid feature names, but DecisionTreeClassifier was fitted with feature names
  warnings.warn(
```



```
C:\Users\Admin\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\utils\validation.py:2739: UserWarning: X does not have valid feature names, but RandomForestClassifier was fitted with feature names
  warnings.warn(
```



## Conclusion

All four models (SVM, k-NN, Decision Tree, and Random Forest) achieved 100% accuracy on the test set. This suggests that the Iris dataset is well-separated and easily classified using these models.

SVM performed well, likely due to the clear decision boundaries between species. k-NN worked perfectly, benefiting from the small dataset size. Decision Tree & Random Forest both achieved perfect classification, showing that tree-based methods are effective.